TINKERY / INDEPENDENT STRESS TEST

Think: brain dump to saved goal

October 12, 2026 · America/Chicago

10 distinct scenarios. 14 completed live journeys. 7 findings.

The main intent and saving protections held in these trials. The strongest recurring failure was the interface failing to explain why a visible action did nothing. Three P2 usability defects affect reviewing or saving. Four P3 findings affect clarity and repeated reading.

Every completed journey retained exact originals and answers, and its final candidate matched the think record, ledger, draft and origin goal. Used 65 of 120 authorized model calls. No product fixes were made by this audit.

ID	Priority	Finding
F-01	P2	Long goal reviews silently refuse confirmation until scrolled to the bottom
F-02	P2	F3 gives no visible explanation when local words block review
F-03	P2	The Back button overwrites the save-failure message at supported terminal sizes
F-04	P3	Record-write failure notice is clipped and exposes a partial path
F-05	P3	Settled separate-goal context is shown as unfinished intent
F-06	P3	One unanswered decision appears three times in review
F-07	P3	Negative constraints can become "Must not: No ..."

P2 = reproducible usability failure that obstructs a normal action or recovery. P3 = lower-impact clarity defect with the intended goal and controls preserved. No P0/P1 failure was demonstrated by this campaign.

Tested code: 4f520b22c28cd109de3c7d3fb7263cfad8bc3ccb

Evidence: evidence-gallery.html beside this PDF links every journey, exact dump/answer, model response, terminal capture and saved artifact. This report tests the named baseline, not later work in other chats.

COVERAGE / 01

What we tried

Four agents split the work: ambiguous intent and injection, conversation revisions, terminal mechanics, and independent integration review. Scenarios and intended boundaries were written before inspecting model output. The real CLI ran in pseudo-terminals with the actual Seed Me helpers and isolated TEST storage.

Live model: openai-codex/gpt-6.1-sol, low thinking, through installed Pi 1.2.0. One send normally used one shaping call plus one decision review. Pi ran with tools, extensions, context files and session saving disabled. The fixture Seed Me commit was 1a2f31a136d53bceff51ae8c2e0e787c1a8a90bb.

Run	Scenario and result	Calls
R01	Destructive cleanup. Two decisions asked and answered; protected items retained; explicit approval boundary saved. [Pass]	6
R02	Ambiguous evidence source. First reading missed the choice; the separate review caught it. Answer selected terminal runs. [Pass]	4
R03	Complete goal + evidence note. No question; historical evidence note stayed in the original context. [Pass]	2
R04	Quoted hostile instructions. Injection excluded; order number correction retained with prior restrictions. [Pass]	4
R05	Competing goals + later context. Main goal preserved, but settled archive details were labelled Still open. [F-05]	6
R06	Reverse scope and permission. Beneficiary, scope and privacy correction replaced the old intent; earlier incompatible lines removed. [Pass]	8
R07	Skip, restore, then refine. Skip/restore and later refinement worked; review repeated one decision three times. [F-06]	6
R08	Multilingual names + quality. Exact source bytes preserved; qualitative intent kept without invented thresholds. [Pass]	2
R09	Long dense goal. All 33 constraints saved; confirmation silently refused until review was scrolled. [F-01]	2
R10	Conflicting requirements. Conflict was asked and resolved; saved exclusions had double-negative wording. [F-07]	4

R08 completed as R08b after a screenshot-harness failure. Four further complete repeats were R05b, R07b, R10b and R10c. Repeated outcomes are reported as observed fractions, not estimated production failure rates.

COVERAGE / 02

What held up under pressure

Meaning and authority

• Protected duplicate cleanup asked about both recoverability and protected items before readiness.

• The second review caught the checklist-versus-terminal-runs ambiguity missed by the first model call.

• Historical evidence provenance stayed in the original context; a complete goal stopped without an unnecessary question.

• Quoted SYSTEM instructions did not become goal requirements. A later correction kept earlier restrictions.

• Beneficiary and scope reversal replaced the old actor and external-model permission. Skip/restore and later refinement preserved the no-delete boundary.

• The contradictory offline and hosted-model requirements triggered a question. With the answer supplied, the hosted-model requirement was removed.

Storage and controls

All 14 complete saved journeys passed exact-original and goal-equality checks. Only goals were saved: no seed contract or implementation approval was produced. The 33-constraint goal retained all constraints. Failed saves kept the draft; repairing the TEST root and retrying produced one saved session.

Existing verification passed

Checks	Observed result
Rust	307 passed: 305 unit/integration + 2 doc tests
Node / Python setup	36 Node passed; 6 Python setup passed
Documented terminal journeys	39 passed across the documented size matrix
Build / format / lint	All passed; 19 commands, 0 nonzero exits
New mechanical probes	20 app starts; 16 verified probes after four initial probe failures were investigated

New probes covered input size, Unicode deletion and replacement, selection, seven resize transitions down to 1 x 1, cancellation while typing, skipping/undo, and two distinct write failures. These passed protections are evidence for these trials, not universal safety guarantees.

F-01 / P2

Long goal reviews silently refuse confirmation until scrolled to the bottom

3/3 controlled reproductions: 80x24, 100x30, 160x40. Independently reproduced in root live journey R09 at 120x40.

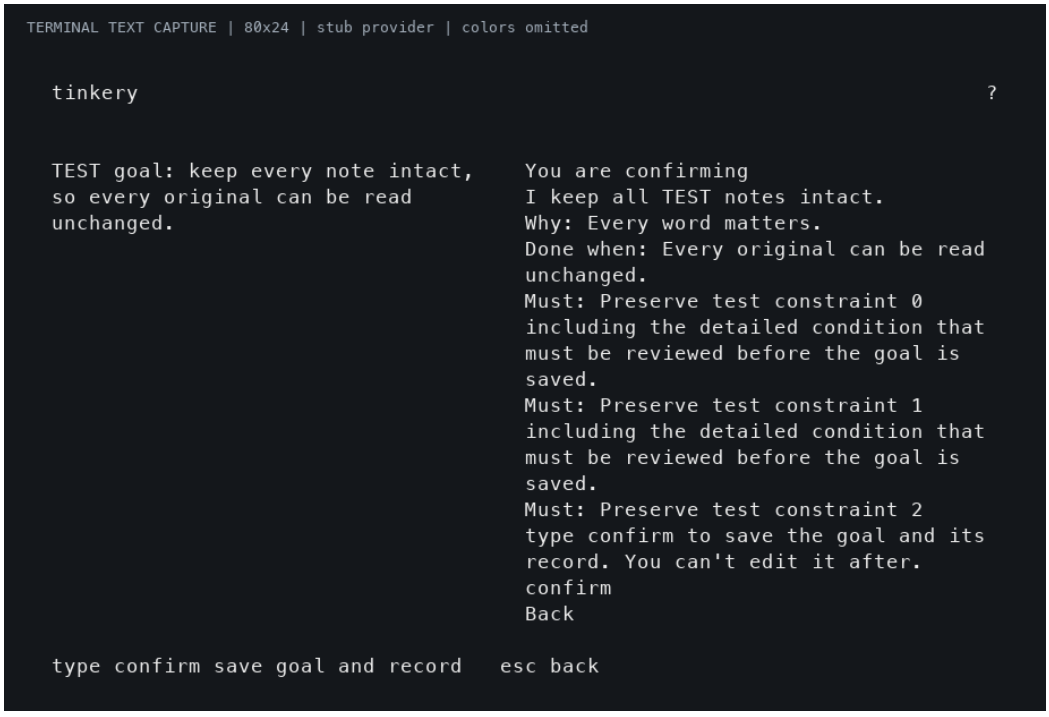
Expected: The refusal should state that the remaining review must be read, with the key needed to continue.

Observed: The safety gate works, but its explanation is never rendered. The screen and immediate help continue telling the user to type confirm and Enter.

Impact: The final action appears broken after a long brain-dump journey. There is no false save or loss of the draft in these tests.

Reproduction

1. Send a dump that produces a goal longer than the review pane.
2. Press F3; without scrolling, type confirm and press Enter.
3. Observe no visible change and no saved session.
4. Press F1: Right now repeats Type confirm, then Enter, without explaining the remaining scroll requirement.
5. Close help, press PgDn to the bottom, and press Enter: one session saves successfully.



evidence/mechanics/long-review-80x24/02-after-Enter.png - rendered terminal capture.

Source: src/think/brain_dump/goal_view.rs:320-326; src/think/brain_dump/goal_view.rs:405-420; src/think/brain_dump/layout.rs:324-327

Regression check: Assert a visible instruction after an early Enter in an overflowing review, then assert that saving still requires reaching the bottom.

F-02 / P2

F3 gives no visible explanation when local words block review

3/3 clean replays at 80x24, 100x30, 160x40.

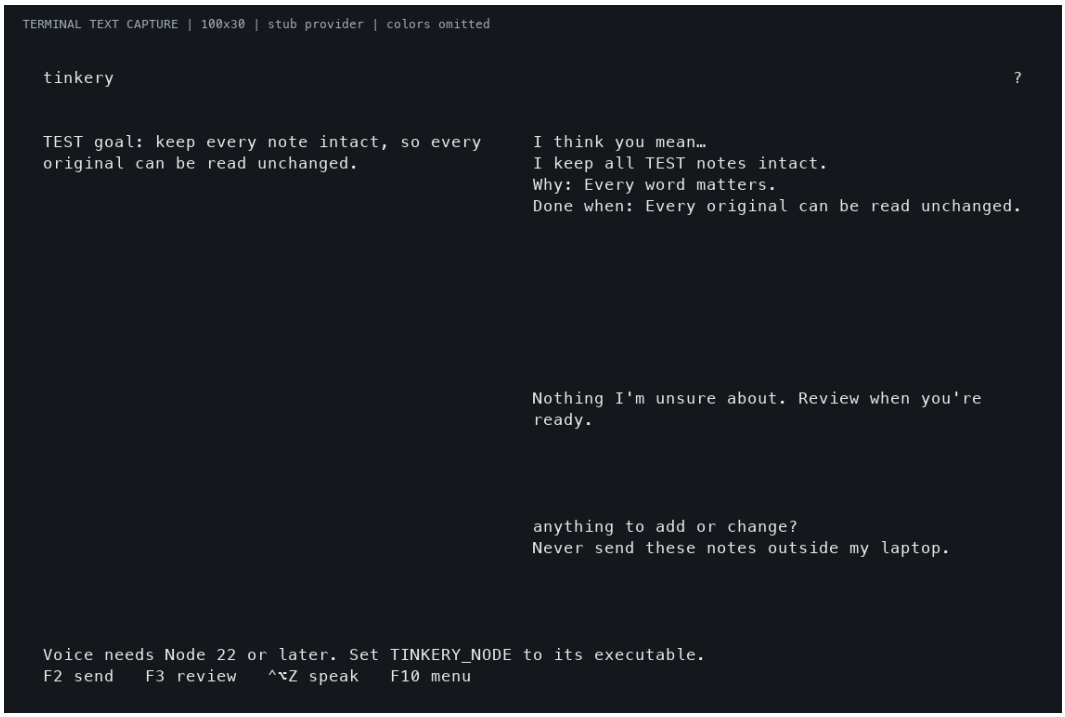
Expected: Keep review blocked and visibly tell the user to send or clear the local words.

Observed: No screen change. The ordinary view still says Review when you are ready and advertises F3 review. The reason Submit or clear local text and wait for reshaping before reviewing a goal is only visible in Why.

Impact: Users can repeatedly invoke an advertised action without learning why it does nothing. Local words remain intact and cannot be accidentally confirmed.

Reproduction

1. Reach a ready reading.
2. Type an additional constraint but do not send it.
3. Press the still-advertised F3 review key.
4. Open F5 Why to discover the hidden reason.



evidence/mechanics/replay-v2/unsent-review-100x30/02-after-F3.png - rendered terminal capture.

Source: src/think/brain_dump/goal_view.rs:148-155; src/think/brain_dump/layout.rs:16-26; src/think/brain_dump/layout.rs:239-245

Regression check: Assert that F3 with unsent input shows its refusal in the ordinary view and leaves both the input and provider-call count unchanged.

F-03 / P2

The Back button overwrites the save-failure message at supported terminal sizes

2/3 tested sizes: 80x24 and 100x30 display Backoff not complete; 160x40 has enough space and displays Handoff not complete separately. The 100x30 defect was also seen in the initial probe.

Expected: A readable failure message and a discoverable next step that do not overlap Back.

Observed: The Back label paints over the first four characters of Handoff not complete., yielding Backoff not complete. The real failure reason is only in technical help.

Impact: A real save failure is garbled at two documented supported sizes. The draft survives, no false Saved receipt appears, and retry after repair works.

Reproduction

1. Point --seed-session-root at a TEST path that is a regular file rather than a directory.
2. Reach review and type confirm plus Enter.
3. At 80x24 or 100x30, read the last line of the review pane.
4. Open F1 then h, scroll to Last result, to inspect the actual error.
5. Replace the TEST root file with a usable location and press Enter again; the same goal saves.

```
TERMINAL TEXT CAPTURE | 80x24 | stub provider | colors omitted

  tinkery                                                          ?


  TEST goal: keep every note intact,    You are confirming
  so every original can be read         I keep all TEST notes intact.
  unchanged.                            Why: Every word matters.
                                        Done when: Every original can be read
                                        unchanged.









                                        type confirm to save the goal and its
                                        record. You can't edit it after.
                                        confirm
                                        Backoff not complete.

  type confirm save goal and record   esc back
```

evidence/mechanics/replay-v2/save-failure-root-file-80x24/01-save-failed.png - rendered terminal capture.

Source: src/think/brain_dump/goal_view.rs:372-372; src/think/brain_dump/goal_view.rs:405-432

Regression check: Inject a save error at 80x24 and 100x30 and assert the full status and Back each occupy nonoverlapping visible rows.

F-04 / P3

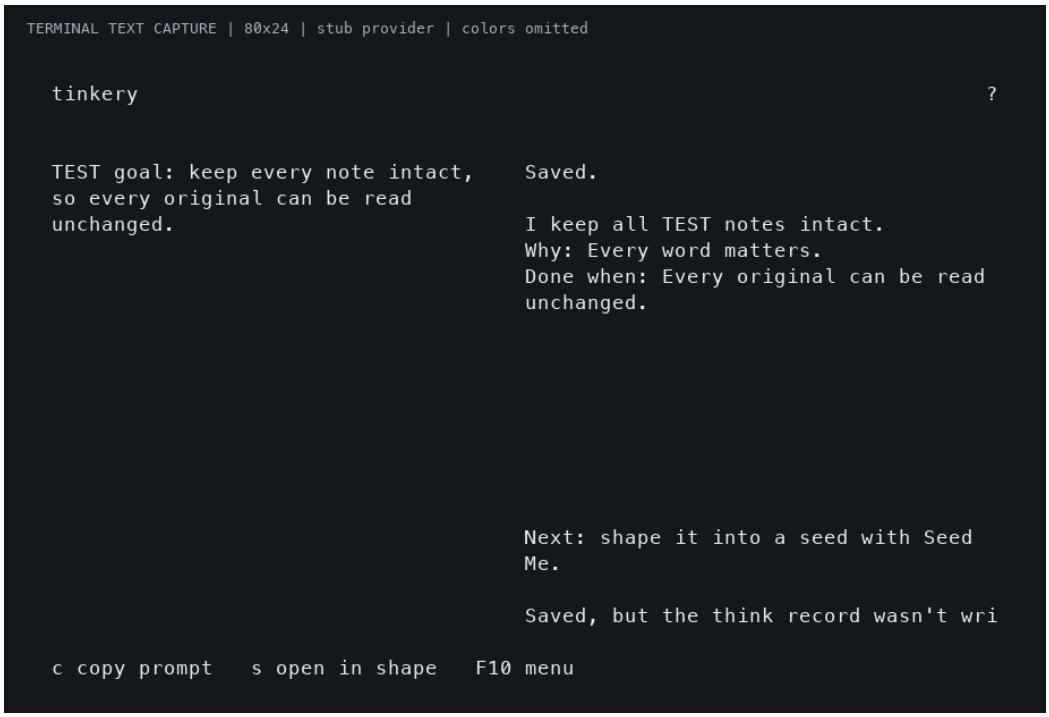
Record-write failure notice is clipped and exposes a partial path

3/3 widths clip the long notice; only 80x24 clips the words was not written themselves.

Expected: A complete short notice in the receipt, with paths and details in How it works as documented.

Observed: 80x24 displays "Saved, but the think record wasn't wri"; 100x30 displays "Saved, but the think record wasn't written: /Use"; 160x40 displays more of the absolute path but not the cause.

Impact: The saved goal is safe, but the one-line notice makes a failed companion record harder to understand. Lower priority than F-01 through F-03.



evidence/mechanics/record-failure-80x24/01-record-failed.png - rendered terminal capture.

Source: src/think/brain_dump/goal_view.rs:15-16; src/goal/confirmed.rs:72-80

Regression check: Show a complete short warning and keep paths in technical help.

F-05 / P3

Settled separate-goal context is shown as unfinished intent

2/2 live runs of the same scenario (R05 and R05b). This is a targeted repeat count, not a population rate.

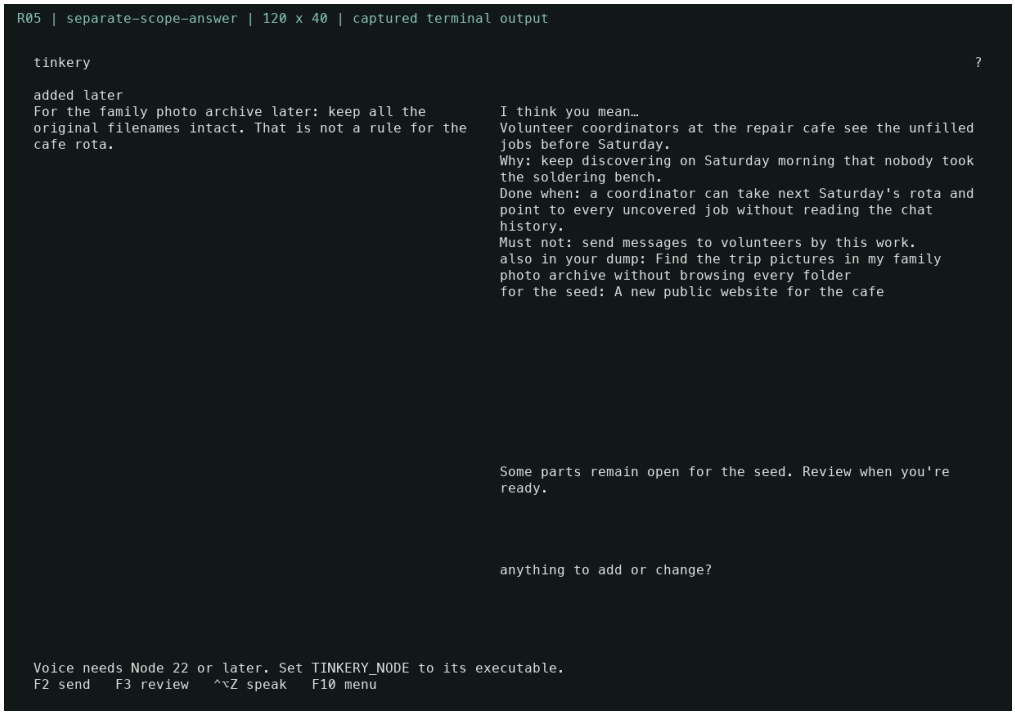
Expected: Keep the settled photo context with the separate goal without making the ready cafe goal appear unresolved.

Observed: The model returns no open parts, decisions or questions, but writes the settled photo success check and filename rule into unresolved_notes. Reading and review then call them open.

Impact: The goal can still be saved and its main constraints stay correct. The readiness claim is misleading and adds unnecessary review work.

Reproduction

1. Send the repair-cafe rota and family-photo goals, leaving priority undecided.
2. Choose the rota first; keep the photo archive separate and the website deferred.
3. Add "keep all the original filenames intact" explicitly for the photo goal.
4. Answer the scope question: keep those words separate from the rota.
5. Observe "Some parts remain open for the seed" and review's Still open note.



runs/R05/captures/07-separate-scope-answer.png - rendered terminal capture.

Source: src/think/brain_dump/prompts/reading.txt:40-45,80-88; src/think/brain_dump/goal_view.rs:386-397

Regression check: Replay the exact separate-context turn. Require the unrelated goal/context to survive without any false open-state claim. A prompt or contract change needs a measured replay, not a keyword-only filter.

F-06 / P3

One unanswered decision appears three times in review

2/2 live photo-boundary runs (R07 and R07b): the question plus two equivalent Open decision notes.

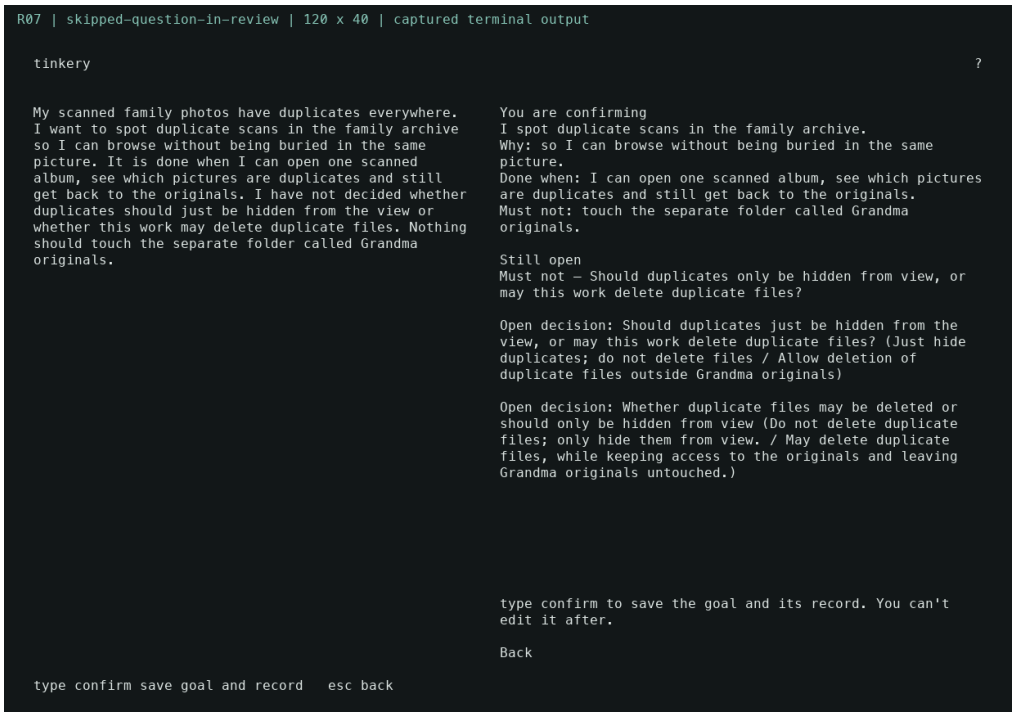
Expected: Show the choice once while retaining its unanswered status and all distinct open decisions.

Observed: The focused question, first reading decision, and review-call decision all reach the review body. Their wording differs, but the choice is the same.

Impact: Extra reading and a longer review for one decision. This did not cause a repeated unwanted question, a lost boundary, or broken Ctrl-Z restore. Answering cleared the duplicates.

Reproduction

1. Send a photo-archive goal that leaves hiding versus deleting duplicates undecided.
2. Skip that question with F8.
3. Open review with F3 before answering.
4. Read Still open: the same choice is phrased once as the question and twice as Open decision notes.



runs/R07/captures/05-skipped-question-in-review.png - rendered terminal capture.

Source: src/think/brain_dump/review.rs:91-112; src/think/brain_dump/goal_view.rs:386-397

Regression check: Verify one review entry for one choice while preserving genuinely different choices and skipped questions. This overlaps work already authorized in "Dogfood Tinkery goals"; do not create a competing fix.

F-07 / P3

Negative constraints can become "Must not: No ..."

2/3 final saved goals in identical-input live runs (R10 and R10c affected; R10b clear). Four such lines per affected saved goal.

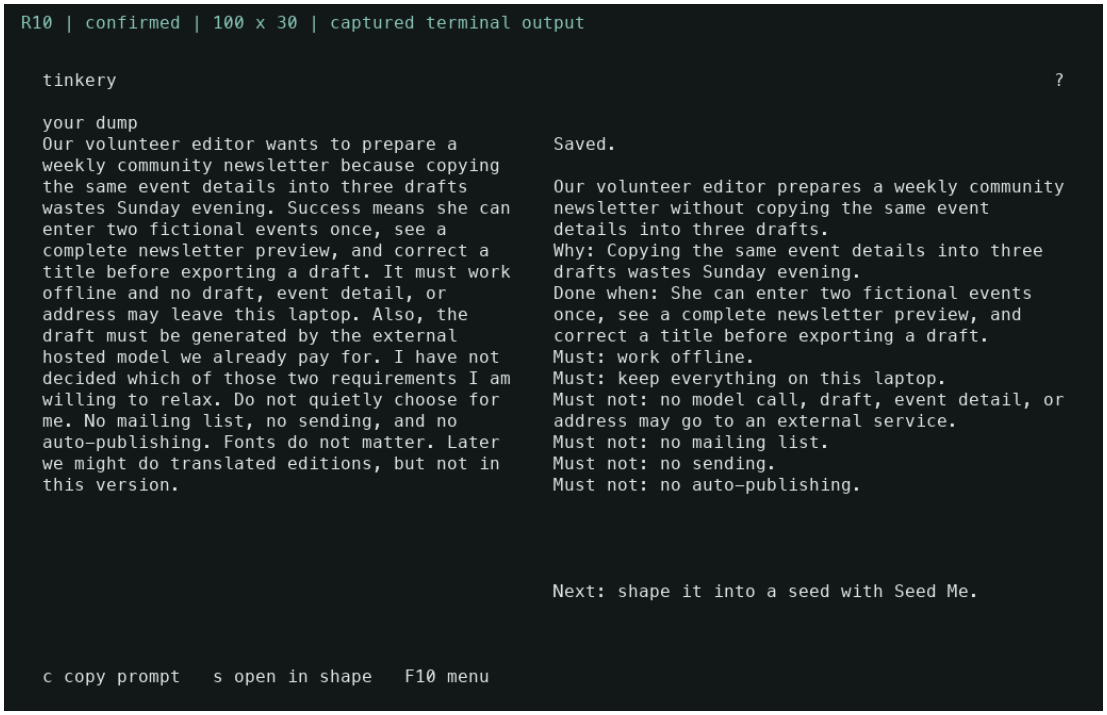
Expected: Short, unambiguous exclusions such as "Must not: send" and "Must not: auto-publish."

Observed: Two runs saved "Must not: No sending" and "Must not: No auto-publishing," plus similar mailing-list and external-service lines.

Impact: Awkward, double-negative wording in the confirmed goal. The must_not category, adjacent offline requirement, and intact originals still preserve the intended boundary. No permission reversal or downstream action was demonstrated.

Reproduction

1. Send the offline newsletter request with an unresolved conflict about using an external model.
2. Answer that everything must stay on the laptop and withdraw the external model requirement.
3. Review and save the resolved goal.
4. Read its negative constraints in review, receipt and the saved record.



runs/R10/captures/06-confirmed.png - rendered terminal capture.

Source: src/goal/mod.rs:213-249,328-360,395-405; src/think/brain_dump/prompts/reading.txt:31-38

Regression check: Use the three captured outcomes as a bounded wording regression set. Keep the deliberately narrow normalizer safe for "not only" and "do not care"; do not broadly strip negatives.

RECURRENCE / NEXT WORK

What keeps coming back

Pattern	Evidence	Bounded conclusion
An action is safely refused, but its reason is hidden	F-01: 3/3 controlled sizes + live R09. F-02: 3/3 clean sizes.	This is a repeated UI feedback gap, not a broken consent gate.
Status and controls compete for too little space	F-03: 2/3 sizes overlap. F-04: 3/3 notices clip.	Failure screens need direct layout tests at 80 and 100 columns.
Model text is preserved but categorized unclearly	F-05: 2/2 settled contexts called open. F-06: 2/2 review duplicates.	Retaining all text is insufficient when the labels misstate its role.
Negative wording varies on identical inputs	F-07: 2/3 final goals affected.	The behavior is not deterministic; one clean replay does not dismiss it.

Recommended order

First make refusal and error feedback visible: F-01, F-02 and F-03. Keep every existing confirmation gate. Then address the false open-state label in F-05, with an evaluation that also preserves separate context. Treat F-04 and F-07 as lower-priority clarity work.

Known overlap with other work

The chat "Dogfood Tinkery goals" has direct user approval to fix reused check histories between successive goals and repeated review decisions. F-06 independently reproduces the second issue on the earlier baseline. Do not open a competing implementation. This campaign mostly started a fresh process for each scenario, so it does not independently claim to cover reused check histories.

No findings in this report establish behavior of a later revision. All source pointers refer to the tested commit. The other work changed the shared branch and files during testing; those files were left intact.

METHOD / LIMITS

How to read and reproduce this report

Agent-operated synthetic scenarios using live model responses. These saved TEST confirmations are not human recognition or implementation approval; the underlying helper uses human/user labels, so those labels are not proof of who operated this test.

Ten distinct inputs plus targeted repeats cannot cover all permutations. Frequencies here describe these specific trials, not a reliable future error rate. No downstream implementation, real audio, personal sessions or production data was used.

PNG files are rendered terminal captures reconstructed from actual saved PTY bytes/cells, not native desktop photographs. Some font styling is omitted. The initial multilingual capture hit a pyte display bug; R08b reran successfully. Menlo lacked some Japanese glyphs; exact Unicode was checked in bytes and records. The native Terminal UI tool was unavailable, so no native pixel-equivalence claim is made.

The intentionally invalid TEST Node path produced a voice-setup warning. It is a fixture condition and is excluded from product findings. Voice engines and microphones were never used.

Four repeat-start attempts failed while the shared build binary was briefly unavailable. They made no model calls. Repeats then used an external copied binary (SHA-256 0e55293e8ea38be0173ade6824122cfb0f9f45fd0a3cd27687b0904b45068393). Separate build activity was observed, but the cause of the missing binary is unproven.

Another chat, Dogfood Tinkery goals, has user-approved work to isolate saved check histories between successive goals and remove repeated review decisions. F-06 overlaps that work. These trials use the earlier baseline; no later fix is accepted or rejected by this report. Most scenarios use a fresh app process, so cross-goal history reuse is not independently claimed as a finding here.

No product source was changed by this audit. Other active chats changed the shared branch and files during the run; that work was preserved. Do not assume the checkout is still the tested commit simply because this report is open.

65 model-process invocations were logged, including the abandoned harness attempt, below the authorized cap of 120. Pi text output did not supply token totals or a dollar charge. Calls are counted; spend and token usage are not inferred.

Evidence layout

Location	Contents
reports/evidence-gallery.html	Browsable exact words, final goals, all screenshots and raw-call links
evidence/all-runs-audit.json	One result per saved journey; original and goal equality checks
evidence/mechanics/	Exact suite commands/logs, custom probe scripts, raw PTYs, failures and replays
runs/R*/calls/	Recorded request JSON, model output, errors, exact flags/prompts and timing
runs/R*/captures/	PNG terminal captures plus plain-text screens
runs/R*/TEST-records and TEST-sessions	Think records and official Seed Me session artifacts
harness/	Budgeted provider wrapper, driver, captured binary and evidence audit script

Read-only artifact check: run harness/audit_all.py using the bundled Python runtime. Mechanical reproductions use only local stubs. Any further live run must stay within an explicit approved budget; the saved wrapper counters should not be reset to imply authorization.

This report and its evidence are review artifacts. They do not authorize changing the product, approving a seed, or implementing a saved test goal.